

# Designing zero knowledge circuits

Daira Hopwood

[daira@jacaranda.org](mailto:daira@jacaranda.org)

@feministPLT

<https://github.com/daira>

@daira on chat.zcashcommunity.com

# Zero knowledge proving systems

- A *statement* is a proposition we want to prove. It depends on:
  - *Instance variables*, which are public.
  - *Witness variables*, which are private.
- Given the instance variables, we can find a short *proof* that we *know* witness variables that make the statement true, without revealing any other information.
- A proof of *knowledge* is stronger and more useful than just proving the statement is *true*. For instance, it allows me to prove that *I know* a secret key, rather than just that *it exists*.
- The proof can be just a string; anyone can verify it without interacting with the prover.
- I'm glossing over some details, such as *setup* and variations of the security properties, which are not the focus of this talk.

# ZK proving systems in the real world

- Since ~2013, zk proving systems have become practical for real-world applications.
- Example: Zcash (<https://z.cash>)
  - Private Bitcoin-like cryptocurrency, with hidden amounts, senders and recipients.
  - (Simplified.) “I know the private key that shows I own  $n$ , which is a valid note with nullifier  $nf$  and a value that balances this transaction.”

Ensuring that a nullifier is not repeated prevents double spending.
- But... *only just* practical:
  - e.g. proof for a private payment at the initial launch of Zcash took  $> 40$  seconds (reduced to 2.5 seconds using some of the techniques described later in this talk).
- This talk isn't about Zcash, but it shows the kind of things we want to be able to prove.

# ZK proving systems in the real world

- The current focus is on proving cryptographic protocols are followed correctly.
- What kinds of things are used in cryptography?
  - Hash functions:  $R = H(B)$
  - One-way functions:  $Q = [x] P$
  - Building blocks:  $B$  is a bit string;  $0 \leq x < y$ ;  $P$  is a valid elliptic curve point; arithmetic; boolean logic; conditionals; ...
  - Recursive validation:  $\pi$  is a valid proof for instance  $X$ .

# Languages for statements

- In the future, statements will be written in high-level languages (e.g. Snarky, Zokrates).
- This talk is not about how to express statements in a concrete programming language. For that see:
  - <https://z.cash/blog/bellman-zksnarks-in-rust/> for bellman (Rust library; used by Zcash)
  - <https://o1labs.org/blog/posts/snarky.html> for Snarky (O'Caml embedded DSL; used by Coda)
  - <https://github.com/Zokrates/ZoKrates> for ZoKrates (dedicated language; used in the Ethereum community).
- All of these systems compile to a language called R1CS (Rank 1 Constraint Systems), which is the subject of this talk.

# R1CS

- Even when programming in a higher-level language, it's necessary to know R1CS.
  - just as understanding the machine model is useful for writing efficient code in conventional programming languages...
  - but even more so, because Snarky, Zokrates, etc. expose many details of the R1CS model (and future proof-oriented languages are also likely to do so).
- We call R1CS programs *circuits*. This talk aims to give a flavour of how R1CS circuits are written and optimized.
- Many different proving systems use R1CS, and it's a current focus of standards development, so this knowledge is transferrable between systems.

# R1CS is only a compilation target

- It is untyped (all values are field elements).
  - Always use types in specifications or when using a proving library / DSL.
- It has no direct support for abstraction.
  - It would be difficult to reverse-engineer the intended statement from an R1CS description. Use abstraction mechanisms supported by the proving library / DSL.
- It's very error-prone.
  - Overflow, missing constraints, unhandled exceptional cases, etc. can cause silent fatal security flaws.

# Fields

- We have instance and witness variables. Variables have values in a *field*.
- A field supports addition, subtraction, multiplication and division of elements, with the following laws:
  - *associativity*:  $(a \cdot b) \cdot c = a \cdot (b \cdot c)$  and  $(a + b) + c = a + (b + c)$
  - *commutativity*:  $a \cdot b = b \cdot a$  and  $a + b = b + a$
  - *identities*:  $a + 0 = a$  and  $a \cdot 1 = a$
  - *inverses*:  $a + (-a) = 0$  and  $a \neq 0 \Rightarrow a \cdot (1/a) = 1$  (we write  $a \cdot (1/b)$  as  $a/b$ )
  - *distributivity*:  $a \cdot (b + c) = a \cdot b + a \cdot c$
- Examples: real numbers, complex numbers, integers modulo a prime.
- We use *finite fields* for cryptography, because elements have “short”, exact representations. In this talk, we only need integers modulo a prime:  $\mathbb{F}_p$ .



# Consequences of using $\mathbb{F}_p$

- Field elements are *not* integers.
- We can use them to *represent* integers, if we're careful.
- “Short” means ~255 bits
  - which is enough to represent a lot of integers, but
  - we **always** need to be careful of overflow.
- We can also use them to represent bits: 0 or 1.
  - this is inefficient (but often necessary)
  - because the proving system still operates on the full field width.
- We can use them to represent themselves!
  - very useful for elliptic curve cryptography
  - but only if the field matches the prime we need.

# Rank 1 constraints

- A *rank 1 constraint system* is a set of *rank 1 constraints*, each of the form:

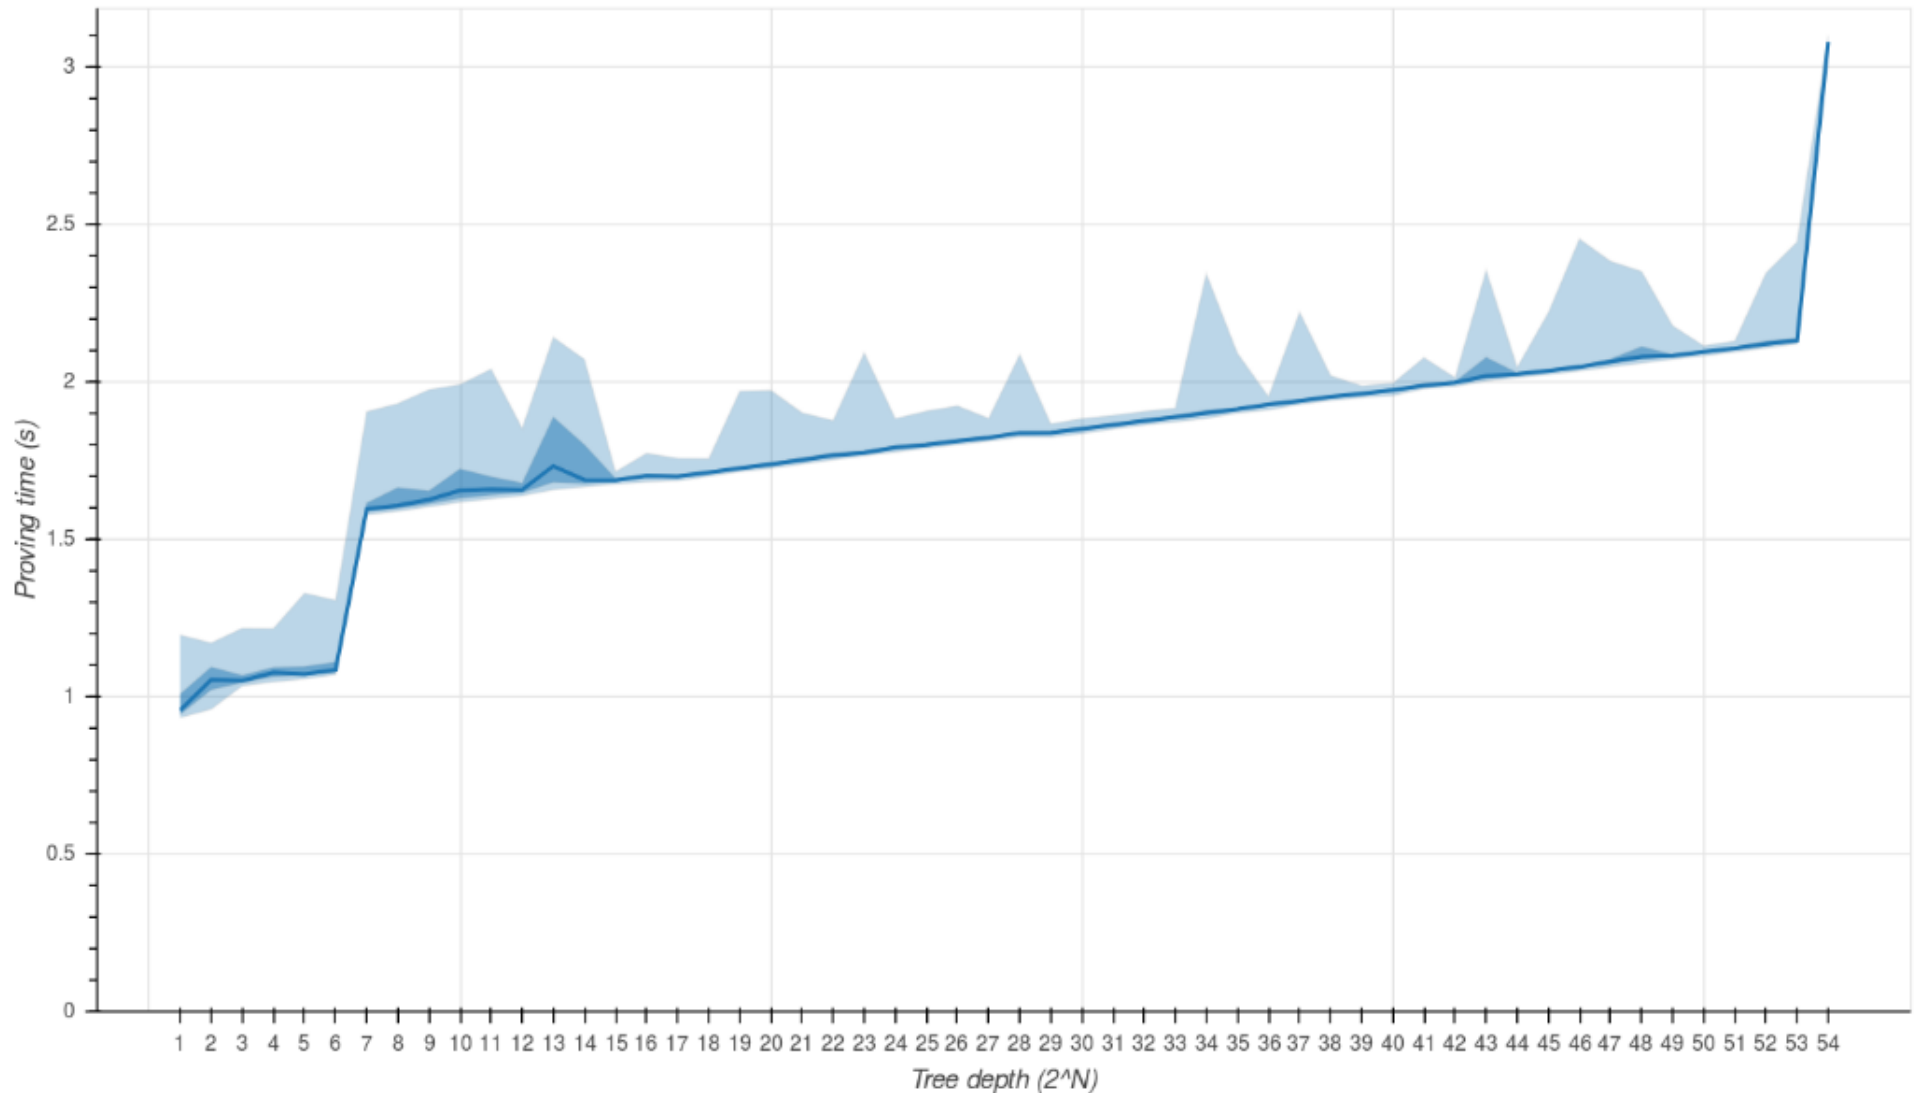
$$(A) \times (B) = (C)$$

where  $(A)$ ,  $(B)$ ,  $(C)$  are each *linear combinations*  $c_1 \cdot v_1 + c_2 \cdot v_2 + \dots$

- The  $c_i$  are *constant* field elements, and the  $v_i$  are instance or witness variables (or 1).
- By convention we write “ $\times$ ” for the multiplications that we need to count, but it’s the same operation as “ $\cdot$ ”.
- Think of general multiplications and divisions as costing 1 constraint; additions, subtractions and scaling by constants are “free”.
  - This is not quite accurate but still a good mental model for optimization. The cost of the circuit will be roughly dependent on the number of constraints (*not* the complexity of the linear combinations).

# Dependence of proving time on circuit size

Sapling input circuit performance



# Rank 1 constraints

- R1CS is a *constraint language*.
- The inputs and outputs of a subcircuit are not predetermined.
  - $(A) \times (B) = (C)$  doesn't mean  $(C)$  is computed from  $(A)$  and  $(B)$ , just that  $(A)$ ,  $(B)$  and  $(C)$  are consistent.
  - more generally, an implementation of “ $x = f(a, b)$ ” doesn't mean that  $x$  is computed from  $a$  and  $b$ , just that  $x$ ,  $a$ , and  $b$  are consistent.
- Constraint languages can be viewed as a generalization of functional languages:
  - everything is referentially transparent and side-effect free
  - there is no ordering of constraints
  - composing two R1CS programs just means that their constraints are simultaneously satisfied.

# Correctness and efficiency

- Multiplication and linear combinations allow us to represent arbitrary circuits:
  - $(1 - b) \times (b) = 0$  is a *boolean constraint* for  $b$ .
  - “ $a$  AND  $b$ ” can be implemented as  $(a) \times (b)$ , and “NOT  $b$ ” as  $1 - b$ .
  - This is a complete set of boolean/bit operations.
- The question is how to represent circuits
  - efficiently (roughly: in the fewest constraints), and
  - correctly (expressing what we intended).
- Correctness is a prerequisite for security. It is not sufficient (we also need to be implementing a secure protocol), but it is necessary.
- Efficient use of fields can allow *4 orders of magnitude* improvement over naive use of bit operations.
  - multiplying two 255-bit numbers would require ~34000 bit operations, but we can do it in one constraint.

# Starting with the basics

- We've already seen  $(1 - b) \times (b) = 0$ .
- This is an instance of a common pattern:  
 $(A) \times (B) = 0$  implements “ $A = 0$  or  $B = 0$ ”.
- We can substitute  $A = P - Q$  and  $B = R - S$ , to get “ $P = Q$  or  $R = S$ ”.
- We did not “reify”  $A = 0$  and  $B = 0$  as boolean variables and explicitly implement OR.
  - Don't reify constraints as booleans unless you have to. There is a way to do that, but it's complicated (and costs 3 constraints).
- This isn't the only way to do a boolean constraint:  $(b) \times (b) = (b)$  also works.
  - It's useful to be able to recognise alternative ways of doing the same thing when reverse engineering R1CS circuits written by others / generated by other libraries.

# Inequalities

- What about  $A \neq 0$  ?
- $0$  is the only field element that doesn't have a multiplicative inverse ( $1/0$  does not exist).
- So  $(A) \times (A_{inv}) = 1$  ensures that  $A \neq 0$ .
- We've added a witness variable,  $A_{inv}$ , that is just an implementation detail rather than part of the original statement. This is very common.
  - It is a witness variable, not an instance variable, because it would leak information, in this case the value of  $A$ , if made public.
- $A - B \neq 0$  is equivalent to  $A \neq B$ . From now on we'll take equivalences like this as obvious.

# Division

- $a = c/b$  is equivalent to  $(a) \times (b) = (c)$ .
- What does  $(a) \times (b) = (c)$  do when  $b = 0$ ?
  - It constrains  $c$  to 0 and leaves  $a$  unconstrained.
  - This makes sense, but is probably not what you want. So don't do that (either constrain or prove  $b \neq 0$ ).
- Notice that division is the same cost as multiplication. This is different from computing inverses in a prime field “outside the circuit”, which is much more expensive than multiplication.
  - Technically, you still need to compute the division when proving. But the cost of that is far outweighed by the per-constraint cost of proving.
  - The different relative costs of operations may lead us to choose different algorithms and representations.



# Inversions

- Division being expressed via multiplication is a special case of a general principle: inversions are easy to express.
- If  $f$  is invertible,  $y = f^{-1}(x)$  is equivalent to  $f(y) = x$ .
- (In the previous slide,  $f^{-1}(x)$  was  $x/b$ .)

# Boolean operations: AND

- Let's use  $n$ -ary AND as an example.
- How many constraints do we need to implement  $b = \text{AND}_{i \in \{1..n\}} a_i$ ?
- There's an obvious implementation in  $n - 1$  constraints. Can we do better?
- We know the answer is boolean:

$$(1 - b) \times (b) = 0$$

- If the answer is  $1$ , then all of the  $a_i$  must be  $1$ :

$$(n - \sum_{i \in \{1..n\}} a_i) \times (b) = 0$$

- If the answer is  $0$ , then not all of the  $a_i$  must be  $1$ :

$$(n - \sum_{i \in \{1..n\}} a_i) \times (inv) = (1 - b)$$

- So, at most 3 constraints independent of  $n$ .
- Notice how we're making use of the representation of booleans as  $0$  or  $1$  and doing arithmetic on them, in order to take advantage of "free" linear combinations. (This won't overflow because  $n < p$ .)
- $n$ -ary OR can be implemented similarly.

# Range constraints

- For binary ranges  $a \in \{0..2^n-1\}$ :  
boolean-constrain  $a_i$  for  $i \in \{0..n-1\}$   
 $(\sum_{i \in \{0..n-1\}} a_i \cdot 2^i) \times (1) = (a)$
- For  $a \in \{0..c-1\}$ , let  $n$  be the bit length of  $c$  and constrain  $a \in \{0..2^n-1\}$  and  $a + 2^n - c \in \{0..2^n-1\}$ .
- There's a more efficient approach that depends on the bit pattern of  $c - 1$ .  
[Zcash specification, appendix A.3.2.2]
- The “ $\times (1)$ ” is technically redundant: we could perform a substitution to eliminate it. In general it's always possible to substitute linear combinations rather than adding another constraint.
  - Your constraint proving library/language should do the substitutions for you. bellman supports this; not sure about Snarky or ZoKrates.

# Boolean operations: XOR

- $c = a \text{ XOR } b$  can be implemented as:

$$(2 \cdot a) \times (b) = (a + b - c)$$

which is equivalent to  $c = a + b - (a \text{ AND } b) \cdot 2 : 0$ .

- What about  $n$ -ary XOR? How many constraints do we need to implement  $b = \text{XOR}_{i \in \{1..n\}} a_i$ ?

- $b$  is the least significant bit of  $\sum_{i \in \{1..n\}} a_i$ .

boolean-constrain  $b_j$  for  $j \in \{0..\text{ceiling}(\lg(n))-1\}$

$$(\sum_{i \in \{1..n\}} a_i) \times (1) = (\sum_{j \in \{0..\text{ceiling}(\lg(n))-1\}} b_j \cdot 2^j)$$

- Now the answer is  $b_0$ .
- So, at most  $\text{ceiling}(\lg(n)) + 1$  constraints independent of  $n$ .

# Conditionals

- A selection constraint  $(b ? x : y) = z$ , where  $b$  has been boolean-constrained, can be implemented as:

$$(b) \times (y - x) = (y - z)$$

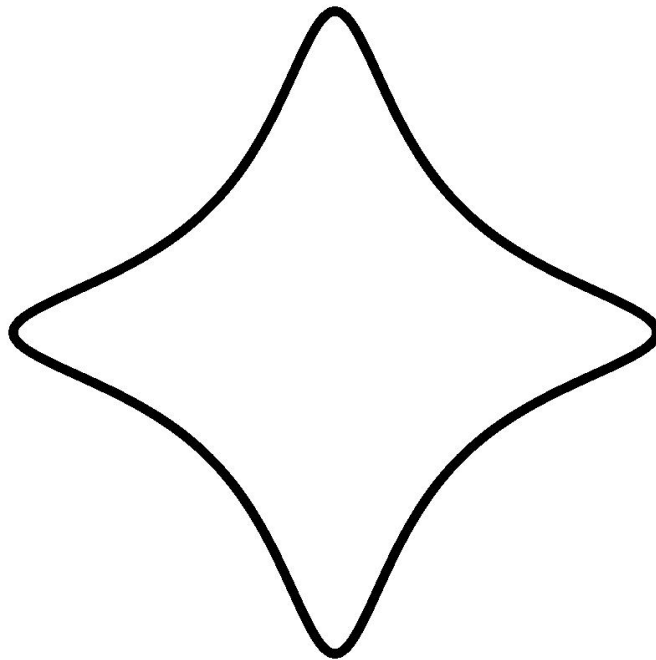
- We can see this is correct by case analysis on  $b$ :
  - If  $b = 1$  then  $y - x = y - z$  therefore  $z = x$ .
  - If  $b = 0$  then  $y - z = 0$  therefore  $z = y$ .

# Elliptic curves

- The most commonly deployed public key cryptosystems use elliptic curves.
- With elliptic curves we can also implement collision-resistant hash functions.
  - Hash functions are the “nails” of cryptography, used everywhere.
- EC crypto can be very efficient in a circuit, compared to symmetric crypto:
  - SHA-256 takes ~27000 constraints.
  - A comparable elliptic curve Pedersen hash takes ~864 constraints, not including boolean-constraining the input.
  - This is because SHA-256 is mainly bit operations, while Pedersen makes full use of the field.
  - This is completely the opposite situation to crypto “outside the circuit”.

# Edwards curves

- Equation of a circle:  $u^2 + v^2 = 1$
- Equation of an Edwards curve:  $a \cdot u^2 + v^2 = 1 + d \cdot u^2 \cdot v^2$
- Over real numbers, the curve looks something like:



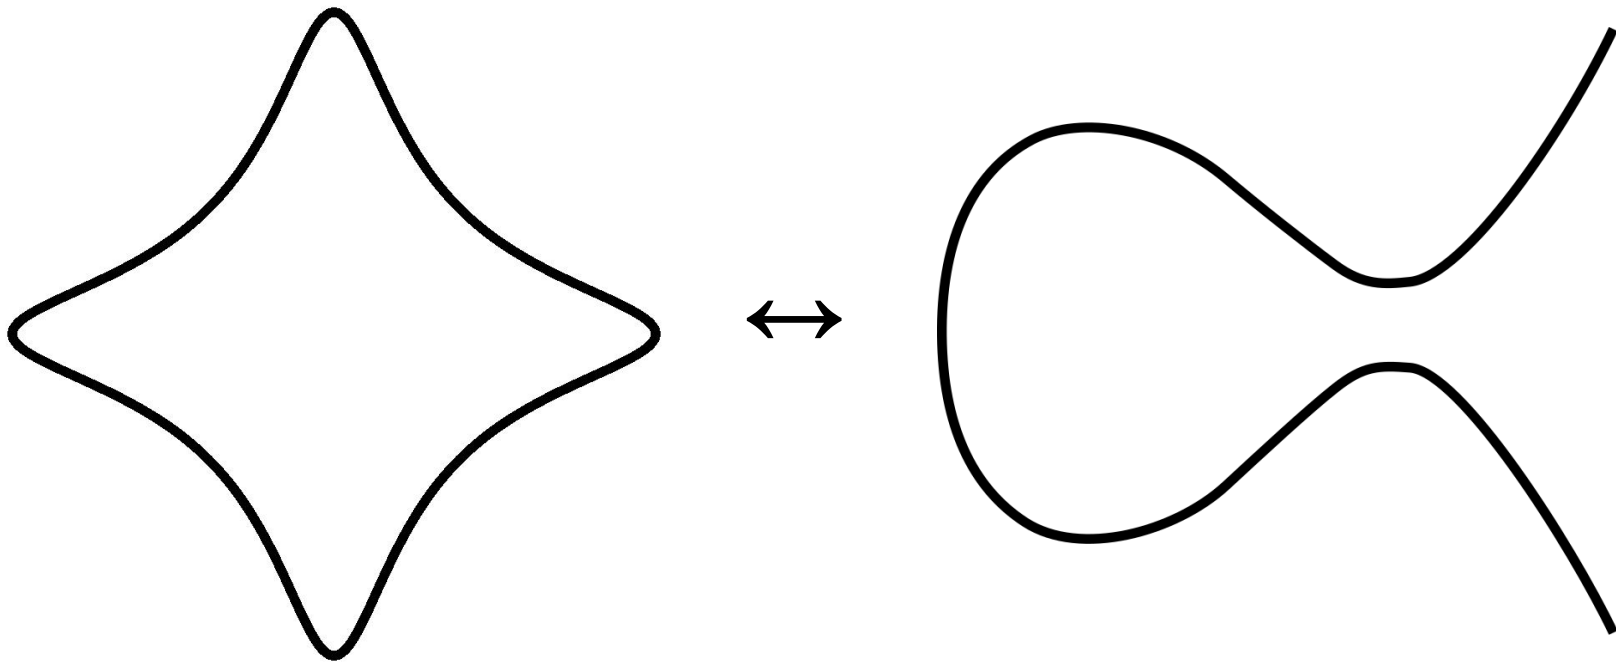
# Edwards arithmetic

- The circuit implementation basically follows the textbook “affine” equations:
  - naive: 7 constraints for addition and doubling
  - optimized: 6 constraints for addition, 5 for doubling
  - Details in the Zcash protocol spec, appendix A.3.3.
- Circuit implementations of elliptic curve arithmetic are actually simpler than out-of-circuit ones, because field division is as efficient as multiplication.



# Montgomery curves

- Each Edwards curve is “birationally equivalent” to a Montgomery curve:



# Montgomery arithmetic

- For a Montgomery curve, addition takes 3 constraints, and doubling takes 4 constraints
- ... but the Montgomery addition doesn't work in all cases; we have to prove that the exceptional cases don't occur.
- We can use the birational equivalence to convert between the fast-but-tricky Montgomery curve, and the slower-but-easier Edwards curve.
- Best to leave this optimization to libraries that are thoroughly reviewed.
- Edwards scalar multiplication  $[x] P$ :
  - fixed  $P$ : 750 constraints, variable  $P$ : 3252 constraints.
- Optimized Montgomery scalar multiplication  $[x] P$ :
  - fixed  $P$ : 506 constraints, variable  $P$ : 2249 constraints (a third better).

# Side rant about correctness proofs

- Common wisdom about use of proofs of (conventional) program correctness: “too hard”, “not ready for prime time”, “the tooling is not there”, “doesn’t scale to real-world programs”, “too hard to maintain when program changes”.
- No! DO PROOFS OR YOU WILL FAIL
  - Optimizations used in Zcash Sapling had proofs of correctness, so this *is* practical.
- Do not whine about needing to do proofs. If you can’t do them, ask a mathematician / cryptographer / appropriate expert. There is a cultural problem with viewing proofs as rocket science, don’t make it worse.
- You don’t necessarily need to use formal theorem provers.
- Do proofs about things that are non-obvious
  - to you, or to a reviewer
  - a lot of things are obvious because the constraint system directly matches the high-level specification.
- Typical proofs are of “this unhandled case can’t occur”, “these algorithms are equivalent”. They will mostly stay valid, or be adapted easily, for changes in the lower-level detail of the constraint system.

# Recursive validation

- Suppose we want to validate a zk proof inside a circuit.
- This allows proving a tree of computation of arbitrary size, in just one proof with constant size and validation time.
- This is a bit of a tour de force and requires much more math than we have time for. The key component is a “pairing”, which is another kind of elliptic curve algorithm.
- Pairings can be implemented in ~7000 constraints, and a full validation (for Groth16) in ~25000 constraints (i.e. less than SHA-256).
  - These are preliminary numbers for a specific curve that might not be quite secure enough, but further optimizations are possible.

# Conclusions

- Writing R1CS programs is interesting and fun...
- but very error-prone. Better tools will be needed.
- Huge performance gains are possible by choosing the right algorithms, representations, and fields.
- Do proofs of correctness where needed.

# Bonus slides: optimization

- Find equivalent expressions of algorithms and use the one with the fewest constraints.
- If expressions are equivalent except for corner cases:
  - constrain the corner cases not to occur, or
  - (better, because no extra constraints) prove that they never occur.
- Switch between multiple representations.
- Change the higher-level protocol to avoid/minimize use of expensive primitives.
- Find non-optimizable things first. Try to reuse values that are unavoidably needed.
- Use algebraic rearrangement to find common subexpressions / make the remaining computations linear.
- Linear expressions are (almost) free. If you are left with linear constraints, remove them by substituting into uses.
  - Ideally, your proving library API / DSL should make this easy.

# Bonus slides: optimization

- Merge to do two things at once.
  - Example: merging with boolean constraints in constant comparisons (Zcash spec A.3.2.2).
- Specialize for constants.
  - Example: lookup from a constant window table in fixed-base scalar multiplication.
- Use nondeterminism.
  - Examples: proving that a value is a square, or non-zero.
- We have concentrated on minimizing number of constraints, but there is also a cost to computing the witness. This can often be optimized by combining operations.
- $n$ -ary operations can often be made less than  $n$  times as expensive as 2-ary.
- Trade operations inside the circuit for operations outside.
- Booleans are (typically) represented as field elements and you can do non-boolean arithmetic on them.
- The most efficient operations are those you can remove.

# Bonus slides: opinionated advice

- Avoid 90s crypto
  - hashes before SHA-256; ciphers before AES.
  - They tend to be inefficient, particularly so in a circuit, even before considering security.
- Many standardized algorithms incur expense that is unnecessary for the small fixed input sizes typically used in circuits.
  - e.g. can use BLAKE2s on a single block directly as a PRF, no need for HMAC/HKDF
  - check with a cryptographer if you are not one.
- Scour the cryptographic literature for cheaper primitives (maybe discarded because they weren't competitive in out-of-circuit computation).
- Use personalization. It's typically free or very cheap, and prevents some chosen-protocol and replay attacks.
- Minimize the primitives used. Circuit programming is difficult and the fewer distinct primitives, the less chance of mistakes and the easier review will be.
- But don't be afraid to specialize if it really helps performance.
- Include redundant checks if they simplify the security analysis and are cheap enough.
- Don't spend time optimizing stuff that makes little difference to overall performance. "Premature optimization is the root of all evil" still applies.
- Set a well-defined "good enough" criterion and stick to it.
- If you don't have imposter syndrome about designing zk circuits in 2018, you're probably doing something wrong.